



Reverse Image Search

Architecture Overview

This project allows users to search for visually similar images by uploading an image. It uses deep learning to extract features from images and stores these features in a vector database. When a user uploads a new image, its features are extracted and compared with those in the database to retrieve the most similar matches. The Visual Search system enables similarity-based retrieval of animal faces from a dataset. A query image is passed through a deep learning model that generates a feature embedding, which is then compared to pre-indexed embeddings to find the most similar images. The architecture consists of *four core components* working together:

1. Feature Extraction Layer (Google Colab)

This is where the image dataset is pre-processed and encoded into numerical feature vectors. Generate compact, discriminative, and semantically meaningful embeddings from input images to represent visual similarity. I utilize a pre-trained Convolutional Neural Network (CNN) as a feature extractor. Fine-tuning on a domain-specific dataset ensures the features capture relevant visual traits (e.g., fur texture, ear shape, eye distance).

- Uses a pre-trained *ResNet50 CNN* to generate a 2048-dimensional embedding for each image.
- L2 normalization ensures cosine similarity works effectively.
- Used augmentations during training to encourage invariance (e.g., random crops, flips, brightness jitter).
- Feature vectors are batched and uploaded to Pinecone along with image metadata (e.g., filename, Cloudinary URL).

This layer is run once initially (or periodically) to prepare the dataset.

2. Vector Database (Pinecone)

Pinecone stores all feature vectors and enables fast similarity searches.

- The index is set with 2048 dimensions and cosine similarity metric.
- Each vector is accompanied by metadata, such as image name and URL, making the results human-readable.
- Pinecone returns the top-k most similar vectors when queried with a new image's features.

This acts as the brain of the search system — optimized for high-speed, approximate nearest neighbor lookups.

3. Flask API

This web service handles search requests from the frontend.

- Accepts base64-encoded images from the frontend and processes them in-memory using Pillow.
- Applies the same ResNet50 feature extraction pipeline as used in Colab.
- Queries Pinecone with the extracted vector.
- Parses the result and returns top matches (URLs, filenames, similarity scores) as JSON.

This is the glue between your client-side UI and the vector database.

4. Frontend (HTML + JavaScript)

This is the user-facing interface where visual search is initiated.

- Users can upload or drag an image to start a search.
- Shows a preview of the uploaded image and enables the search button.
- Sends the image to the backend via a POST request.
- Renders similar image thumbnails with similarity percentages.

Components

Component	Technology	Purpose
Feature Extraction	PyTorch (ResNet50)	Converts image into a vector
Vector DB	Pinecone (Free Tier)	Stores vectors and supports similarity search
API Server	Flask (Python)	Serves search endpoint
Frontend	HTML, JavaScript	Image upload and results display
Hosting	Railway hosting	Deploys backend and frontend
Storage	Cloudinary (Free Tier)	Stores original image files for preview

Codebase

To keep this guide concise, the full source code is broken down below:

- 🧠 **Feature Extraction/Flask API** → [Visual Search - Github](#)
- 💻 **WebApp** → [Visual Search App](#)

